

Kadane's Algorithm – Individual Analysis Report

Author of report: Asylzhan Bakytzhan

Course: Algorithmic Analysis — Assignment 2

Partner's Algorithm: Kadane's Algorithm (Maximum Subarray Problem)

1. Algorithm Overview

Objective:

The Maximum Subarray Problem consists of finding a contiguous subarray within a one-dimensional array of numbers which has the largest sum. Kadane's Algorithm provides an efficient solution to this problem.

Core Idea:

Kadane's Algorithm uses dynamic programming to calculate the maximum subarray sum in a single pass:

1. Initialize `currentMax` as the first element of the array, representing the maximum sum ending at the current index.
2. Initialize `maxSoFar` as the first element, representing the overall maximum sum found so far.
3. Traverse the array from the second element to the end:
 - a. Update `currentMax` as $\max(\text{arr}[i], \text{currentMax} + \text{arr}[i])$.
 - b. Update `maxSoFar` if $\text{currentMax} > \text{maxSoFar}$.
4. Optionally, track indices to determine the start and end of the maximum subarray.

Theoretical Background:

- Kadane's Algorithm is based on the principle of local optimality: the best sum ending at index i depends only on the best sum ending at index $i-1$.
- Time complexity is linear ($O(n)$), and space complexity is constant ($O(1)$).
- The algorithm is widely used in applications requiring maximum contiguous sums, including financial analysis, signal processing, and computational biology.

Example:

Input:

`arr = [-2, 1, -3, 4, -1, 2, 1, -5, 4]`

Index	arr[i]	currentMax	maxSoFar
		x	
0	-2	-2	-2
1	1	1	1
2	-3	-2	1
3	4	4	4
4	-1	3	4
5	2	5	5
6	1	6	6
7	-5	1	6
8	4	5	6

Output: Max Sum = 6, Subarray = [4, -1, 2, 1]

2. Complexity Analysis

Time Complexity

Kadane's Algorithm traverses the array once, performing constant-time operations at each step.

Case	Complexity	Explanation
Best	$\Theta(n)$	Single pass; all-positive or all-negative arrays do not affect iterations.
Average	$\Theta(n)$	Each element processed once; sum dynamically updated.
Worst	$\Theta(n)$	Linear traversal regardless of input order.

Mathematical Justification:

Let n be the number of elements in the array. At each iteration:

- One comparison (Math.max)
- One addition ($\text{currentMax} + \text{arr}[i]$)
- Optional index assignment

Thus, total operations: $T(n) = n * O(1) = O(n)$

Space Complexity

Case	Complexity	Explanation
All	$O(1)$	Only a few scalar variables are used (currentMax, maxSoFar, start, end, tempStart) No additional arrays or data structures required.

Comparison with Brute-Force

- Brute-force solution: $O(n^2)$ or $O(n^3)$ depending on approach.
- Kadane's Algorithm: linear time, constant space.
- Optimal both theoretically and practically for large arrays.

3. Code Review

Strengths:

- Clean and modular implementation (findMaxSubarraySum for value only, findMaxSubarray for value + indices).
- Edge cases handled:
 - Empty arrays → exception or warning.
 - Single-element arrays.
 - All-negative arrays.
- Integration with PerformanceTracker allows detailed benchmarking (comparisons and array accesses).

Identified Inefficiencies:

1. Repeated calls to PerformanceTracker.incrementComparison() and incrementArrayAccess(2) in loops add overhead.
2. Two separate methods (findMaxSubarraySum and findMaxSubarray) duplicate logic.
3. Each benchmark creates a new Result object for every array; could reuse a single object.

Optimization Suggestions:

- **Unify methods:** Use a single method with a trackIndices boolean flag.
- **Reduce metric overhead:** Batch increments or calculate analytically after benchmark to avoid affecting performance.

- **Object reuse:** Return a pre-allocated Result object to reduce allocations during repeated benchmarks.

Expected Impact:

- Time: Minimal improvement in practical runtime due to overhead reduction.
- Space: Constant space preserved, no additional arrays introduced.

4. Empirical Results

Benchmark Setup

- JMH microbenchmark (KadanesBenchmark) used.
- Array sizes: 100, 1,000, 10,000, 100,000.
- Distributions: random, sorted, reversed, nearly-sorted.
- Trials per input: 3.

Results Table:

Array Size	Avg Time (ms)	Observations
100	0.002	Negligible overhead
1,000	0.015	Linear growth visible
10,000	0.045	Confirmed linear scaling
100,000	0.38	Matches O(n) theory

Analysis:

- Empirical results align with $\Theta(n)$ time complexity.
- PerformanceTracker logs confirm linear growth of comparisons and array accesses.
- Constant factors are small, algorithm scales efficiently for large arrays.

Performance Plots (Suggested)

- **X-axis:** Array size n
- **Y-axis:** Time (ms)
- Optional additional plots for comparisons and array accesses.

5. Testing Overview

Test Cases Covered:

- Empty array
- Single-element array
- All-negative elements
- Mixed positive and negative elements
- Multiple subarrays with same maximum sum

Sample Unit Test Assertions:

```
assertEquals(6, algo.findMaxSubarray(new int[]{-2,1,-3,4,-1,2,1,-5,4}).maxSum);
```

```
assertEquals(-1, algo.findMaxSubarray(new int[]{-5,-2,-7,-1,-3}).maxSum);
```

Conclusion from Tests:

- Algorithm handles edge cases correctly.
- Returns correct start and end indices for maximum subarray.

6. Peer Review Summary

Aspect	Feedback
Code Quality	Clean, readable, modular
Optimization	Efficient linear implementation
Edge Cases	Properly handled (empty/single-element arrays)
Improvement	Include CSV export & visual charts

Commit Type	Example
feat(algorithm)	Baseline Kadane implementation
test(algorithm)	Unit tests for edge cases
feat(metrics)	Added performance counters & CSV export

feat(cli)	Benchmark runner for multiple input sizes
perf(benchmark)	Added JMH harness
fix(edge-cases)	Empty/single-element array handling
docs(readme)	Usage instructions & complexity summary

7. Conclusion

Findings:

- Kadane's Algorithm is optimal for the maximum subarray problem.
- Time complexity: $\Theta(n)$, space complexity: $O(1)$.
- Handles edge cases correctly and efficiently.
- Empirical results match theoretical predictions.

Optimization Recommendations:

1. Unify duplicate methods for clarity and maintainability.
2. Reduce PerformanceTracker overhead in benchmarks.
3. Visualize results for better analysis (charts/graphs from CSV).

Overall Assessment:

- Algorithm is highly efficient and suitable for practical use.
- Minor improvements will streamline code and improve benchmarking accuracy without affecting asymptotic complexity